

From Autoencoder to Beta-VAE

Aug 12, 2018 by Lilian Weng [autoencoder](#) [generative-mode](#)

Autocoders are a family of neural network models aiming to learn compressed latent variables from high-dimensional data. Starting from the basic autoencoder model, this post reviews several variations, including denoising, sparse, and contractive autoencoders, and then Variational Autoencoder (VAE) and its modification beta-VAE.

[Updated on 2019-07-18: add a section on [VQ-VAE & VQ-VAE-2](#).]

[Updated on 2019-07-26: add a section on [TD-VAE](#).]

Autoencoder is invented to reconstruct high-dimensional data using a neural network model with a bottleneck layer in the middle (oops, this is probably not true [Variational Autoencoder](#), and we will investigate it in details in later sections). A nice byproduct is dimension reduction: the bottleneck captures a compressed latent encoding. Such a low-dimensional representation can be used as an embedding vector in various applications (i.e. search), help data compression, or reveal the underlying generative factors.

- [Notation](#)
- [Autoencoder](#)
- [Denoising Autoencoder](#)
- [Sparse Autoencoder](#)
- [Contractive Autoencoder](#)
- [VAE: Variational Autoencoder](#)
 - [Loss Function: ELBO](#)
 - [Reparameterization Trick](#)
- [Beta-VAE](#)
- [VQ-VAE and VQ-VAE-2](#)
- [TD-VAE](#)
- [References](#)

Notation

Symbol	Mean
\mathcal{D}	The dataset, $\mathcal{D} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)}\}$, contains n data samples; $ \mathcal{D} = n$.

Symbol	Mean
$\mathbf{x}^{(i)}$	Each data point is a vector of d dimensions, $\mathbf{x}^{(i)} = [x_1^{(i)}, x_2^{(i)}, \dots, x_d^{(i)}]$.
\mathbf{x}	One data sample from the dataset $\mathbf{x} \in \mathcal{D}$.
\mathbf{x}'	The reconstructed version of \mathbf{x} .
$\tilde{\mathbf{x}}$	The corrupted version of \mathbf{x} .
\mathbf{z}	The compressed code learned in the bottleneck layer.
$a_j^{(l)}$	The activation function for the j -th neuron in the l -th hidden layer.
$g_\phi(\cdot)$	The encoding function parameterized by ϕ .
$f_\theta(\cdot)$	The decoding function parameterized by θ .
$q_\phi(\mathbf{z} \mathbf{x})$	Estimated posterior probability function, also known as probabilistic encoder .
$p_\theta(\mathbf{x} \mathbf{z})$	Likelihood of generating true data sample given the latent code, also known as probabilistic decoder .

Autoencoder

Autoencoder is a neural network designed to learn an identity function in an unsupervised way. It reconstructs the original input while compressing the data in the process so as to discover a more efficient and compressed representation. The idea was originated in the 1980s, and later promoted by the seminal paper by Hinton & Salakhutdinov, 2006.

It consists of two networks:

- *Encoder* network: It translates the original high-dimensional input into the latent low-dimensional representation. The input size is larger than the output size.
- *Decoder* network: The decoder network recovers the data from the code, likely with larger and longer output layers.




Fig. 1. Illustration of autoencoder model architecture.

The encoder network essentially accomplishes **dimensionality reduction**, just like how we would use Principal Component Analysis (PCA) or Matrix Factorization (MF) for. In addition, the autoencoder is explicitly optimized for the data reconstruction from the code. A good intermediate representation can capture latent variables, but also benefits a **decompression** process.

The model contains an encoder function $g(\cdot)$ parameterized by ϕ and a decoder function $f(\cdot)$ parameterized by θ . The low-dimensional code learned for input \mathbf{x} in the bottleneck layer is $\mathbf{z} = f_\theta(g_\phi(\mathbf{x}))$ and the reconstructed input $\mathbf{x}' = f_\theta(g_\phi(\mathbf{x}))$.

The parameters (θ, ϕ) are learned together to output a reconstructed data sample same as the original input, $\mathbf{x}' \approx f_\theta(g_\phi(\mathbf{x}))$, or in other words, to learn an identity function. There are various metrics to quantify the difference between two vectors, such as cross entropy when the activation function is sigmoid, or as simple as MSE loss

$$L_{AE}(\theta, \phi) = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}^{(i)} - f_\theta(g_\phi(\mathbf{x}^{(i)})))^2$$

Denoising Autoencoder

Since the autoencoder learns the identity function, we are facing the risk of “overfitting” when there are more network parameters than the number of data points.

To avoid overfitting and improve the robustness of the autoencoder, Vincent et al. (2008) proposed a modification to the basic autoencoder. The input is partially corrupted by adding noise: masking some values of the input vector in a stochastic manner $\tilde{\mathbf{x}} \sim \mathcal{M}_D(\tilde{\mathbf{x}}|\mathbf{x})$. Then the model is trained to recover the original input **Note: Not the corrupt one!**.

$$\begin{aligned} \tilde{\mathbf{x}}^{(i)} &\sim \mathcal{M}_D(\tilde{\mathbf{x}}^{(i)}|\mathbf{x}^{(i)}) \\ L_{DAE}(\theta, \phi) &= \frac{1}{n} \sum_{i=1}^n (\mathbf{x}^{(i)} - f_\theta(g_\phi(\tilde{\mathbf{x}}^{(i)})))^2 \end{aligned}$$

where \mathcal{M}_D defines the mapping from the true data samples to the noisy or corrupted ones.






Fig. 2. Illustration of denoising autoencoder model architecture.

This design is motivated by the fact that humans can easily recognize an object or a scene even view is partially occluded or corrupted. To “repair” the partially destroyed input, the den autoencoder has to discover and capture relationship between dimensions of input in order to missing pieces.

For high dimensional input with high redundancy, like images, the model is likely to depend on evidence gathered from a combination of many input dimensions to recover the denoised version (sounds like attention mechanism, right?) rather than to overfit one dimension. This builds up a good foundation learning robust latent representation

The noise is controlled by a stochastic mapping $M_D(\tilde{\mathbf{x}}|\mathbf{x})$, and it is not specific to a particular type of corruption process (i.e. masking noise, Gaussian noise, salt-and-pepper noise, etc.). Naturally corruption process can be equipped with prior knowledge

In the experiment of the original DAE paper, the noise is applied in this way: a fixed proportion of dimensions are selected at random and their values are forced to 0. Sounds a lot like dropout, right? Well, the denoising autoencoder was proposed in 2008, 4 years before the dropout paper ([Hinton, et al. 2012](#));

Sparse Autoencoder

Sparse Autoencoder applies a “sparse” constraint on the hidden unit activation to avoid overfitting and improve robustness. It forces the model to only have a small number of hidden units being activated at the same time, or in other words, one hidden neuron should be inactive most of the time.

Recall that common activation function include sigmoid, tanh, relu, leaky relu, etc. A neuron is activated when the value is close to 1 and inactivated with a value close to 0.

Let’s say there are s_l neurons in the l -th hidden layer and the activation function for j -th neuron in this layer is labelled as $a_j^{(l)}(\cdot)$, $j = 1, \dots, s_l$. The fraction of activation of this neuron $\hat{\rho}_j$ is expected to be a small number ρ , known as sparsity parameter; a common config is $\rho = 0.05$.

$$\hat{\rho}_j^{(l)} = \frac{1}{n} \sum_{i=1}^n [a_j^{(l)}(\mathbf{x}^{(i)})] \approx \rho$$

This constraint is achieved by adding a penalty term into the loss function. The KL-divergence D_{KL} measures the difference between two Bernoulli distributions, one with $n\rho$ and the other with mean $\hat{\rho}_j^{(l)}$. The hyperparameter β controls how strong the penalty we want to apply on the sparsity

$$\begin{aligned}
L_{\text{SAE}}(\theta) &= L(\theta) + \beta \sum_{l=1}^L \sum_{j=1}^{s_l} D_{\text{KL}}(\rho \parallel \hat{\rho}_j^{(l)}) \\
&= L(\theta) + \beta \sum_{l=1}^L \sum_{j=1}^{s_l} \rho \log \frac{\rho}{\hat{\rho}_j^{(l)}} + (1 - \rho) \log \frac{1 - \rho}{1 - \hat{\rho}_j^{(l)}}
\end{aligned}$$






Fig. 4. The KL divergence between a Bernoulli distribution with mean $\rho = 0.25$ and a Bernoulli distribution with mean $0 \leq \hat{\rho} \leq 1$.

k-Sparse Autoencoder

In *k*-Sparse Autoencoder [Makhzani and Frey, 2013], the sparsity is enforced by only keeping the top highest activations in the bottleneck layer with linear activation function. First we run feedforward through the encoder network to get the compressed $\text{codz} = g(\mathbf{x})$. Sort the values in the code vector \mathbf{z} . Only the *k* largest values are kept while other neurons are set to 0. This can be done in a ReLU layer with an adjustable threshold too. Now we have a sparsified $\text{codz}' = \text{Sparsify}(\mathbf{z})$. Compute the output and the loss from the sparsified $\text{codL} = \|\mathbf{x} - f(\mathbf{z}')\|_2^2$. And, the back-propagation only goes through the top *k* activated hidden units.



(a) $k = 70$



(b) $k = 40$



(c) $k = 25$

looks like a real data point $\mathbf{x}^{(i)}$, we follow these steps

1. First, sample a $\mathbf{z}^{(i)}$ from a prior distribution $p_{\theta^*}(\mathbf{z})$.
2. Then a value $\mathbf{x}^{(i)}$ is generated from a conditional distribution $p_{\theta^*}(\mathbf{x}|\mathbf{z} = \mathbf{z}^{(i)})$.

The optimal parameter θ^* is the one that maximizes the probability of generating real data samples.

$$\theta^* = \arg \max_{\theta} \prod_{i=1}^n p_{\theta}(\mathbf{x}^{(i)})$$

Commonly we use the log probabilities to convert the product on RHS to a sum:

$$\theta^* = \arg \max_{\theta} \sum_{i=1}^n \log p_{\theta}(\mathbf{x}^{(i)})$$

Now let's update the equation to better demonstrate the data generation process so as to involve the latent encoding vector

$$p_{\theta}(\mathbf{x}^{(i)}) = \int p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z}) p_{\theta}(\mathbf{z}) d\mathbf{z}$$

Unfortunately it is not easy to compute $p_{\theta}(\mathbf{x}^{(i)})$ in this way, as it is very expensive to check all the possible values of \mathbf{z} and sum them up. To narrow down the value space to facilitate faster search, we would like to introduce a new approximation function to output what is a likely code given any \mathbf{x} , $q_{\phi}(\mathbf{z}|\mathbf{x})$, parameterized by ϕ .




Fig. 6. The graphical model involved in Variational Autoencoder. Solid lines denote the generative distribution $p_{\theta}(\cdot)$ and dashed lines denote the distribution $q_{\phi}(\mathbf{z}|\mathbf{x})$ to approximate the intractable posterior $p_{\theta}(\mathbf{z}|\mathbf{x})$.

Now the structure looks a lot like an autoencoder.

- The conditional probability $p_{\theta}(\mathbf{x}|\mathbf{z})$ defines a generative model, similar to the decoding $f_{\theta}(\mathbf{x}|\mathbf{z})$ introduced above. $p_{\theta}(\mathbf{x}|\mathbf{z})$ is also known as *probabilistic decoder*.
- The approximation function $q_{\phi}(\mathbf{z}|\mathbf{x})$ is the *probabilistic encoder*, playing a similar role as $g_{\phi}(\mathbf{z}|\mathbf{x})$ above.

Loss Function: ELBO

The estimated posterior $q_\phi(\mathbf{z}|\mathbf{x})$ should be very close to the real one $p_\theta(\mathbf{z}|\mathbf{x})$. We can use [Kullback-Leibler divergence](#) to quantify the distance between these two distributions. KL divergence $D_{\text{KL}}(X\|Y)$ measures how much information is lost if the distribution Y is used to represent X.

In our case we want to minimize $D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x})\|p_\theta(\mathbf{z}|\mathbf{x}))$ with respect to ϕ .

But why use $D_{\text{KL}}(q_\phi\|p_\theta)$ (reversed KL) instead of $D_{\text{KL}}(p_\theta\|q_\phi)$ (forward KL)? Eric Jang has a great explanation in his [post](#) on Bayesian Variational methods. As a quick recap,




Fig. 7. Forward and reversed KL divergence have different demands on how to match two distributions.
(Image source: blog.evjang.com/2016/08/variational-bayes.htm)

- Forward KL divergence: $D_{\text{KL}}(P\|Q) = \mathbb{E}_{z\sim P(z)} \log \frac{P(z)}{Q(z)}$; we have to ensure that $Q(z)>0$ wherever $P(z)>0$. The optimized variational distribution $q(z)$ has to cover over the entire $P(z)$.
- Reversed KL divergence: $D_{\text{KL}}(Q\|P) = \mathbb{E}_{z\sim Q(z)} \log \frac{Q(z)}{P(z)}$; minimizing the reversed KL divergence squeezes the $Q(z)$ under $P(z)$.

Let's now expand the equation

The expectation term in the loss function invokes generating samples $\mathbf{fz} \sim q_\phi(\mathbf{z}|\mathbf{x})$. Sampling is a stochastic process and therefore we cannot backpropagate the gradient. To make it trainable reparameterization trick is introduced: It is often possible to express the random var \mathbf{z} as a deterministic variable $\mathbf{z} = \mathcal{T}_\phi(\mathbf{x}, \epsilon)$, where ϵ is an auxiliary independent random variable, and the transformation function \mathcal{T}_ϕ parameterized by ϕ converts ϵ to \mathbf{z} .

For example, a common choice of the form $q_\phi(\mathbf{z}|\mathbf{x})$ is a multivariate Gaussian with a diagonal covariance structure

$$\begin{aligned}\mathbf{z} &\sim q_\phi(\mathbf{z}|\mathbf{x}^{(i)}) = \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}^{(i)}, \boldsymbol{\sigma}^{2(i)} \mathbf{I}) \\ \mathbf{z} &= \boldsymbol{\mu} + \boldsymbol{\sigma} \odot \epsilon, \text{ where } \epsilon \sim \mathcal{N}(0, \mathbf{I})\end{aligned}\quad ; \text{ Reparameterization trick.}$$

where \odot refers to element-wise product





Fig. 8. Illustration of how the reparameterization trick makes the sampling process trainable. (Image source: Slide 12 in Kingma's NIPS 2015 workshop [talk](#))

The reparameterization trick works for other types of distributions too, not only Gaussian. In the multivariate Gaussian case, we make the model trainable by learning the mean and variance of the distribution μ and σ , explicitly using the reparameterization trick, while the stochasticity remains in the random variable $\epsilon \sim \mathcal{N}(0, \mathbf{I})$.






Fig. 9. Illustration of variational autoencoder model with the multivariate Gaussian assumption

Beta-VAE

If each variable in the inferred latent representation \mathbf{z} is only sensitive to one single generative factor and relatively invariant to other factors, we will say this representation is disentangled or factorized. This benefit that often comes with disentangled representations is good *interpretability* and easy generalization to a variety of tasks.

For example, a model trained on photos of human faces might capture the gender, skin color, hair color, hair length, emotion, whether wearing a pair of glasses and many other relatively independent factors in separate dimensions. Such a disentangled representation is very beneficial to facial image generation.

β -VAE ([Higgins et al., 2017](#)) is a modification of Variational Autoencoder with a special emphasis on discovering disentangled latent factors. Following the same incentive in VAE, we want to maximize the probability of generating real data, while keeping the distance between the real and estimated posterior distributions small (say, under a small constant δ):

$$\begin{aligned} & \max_{\phi, \theta} \mathbb{E}_{\mathbf{x} \sim \mathcal{D}} [\mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \log p_\theta(\mathbf{x}|\mathbf{z})] \\ & \text{subject to } D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}) \| p_\theta(\mathbf{z})) < \delta \end{aligned}$$

We can rewrite it as a Lagrangian with a Lagrangian multiplier β under the [KKT condition](#). The above optimization problem with only one inequality constraint is equivalent to maximizing the following equation $\mathcal{F}(\theta, \phi, \beta)$:

$$\begin{aligned} \mathcal{F}(\theta, \phi, \beta) &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \log p_\theta(\mathbf{x}|\mathbf{z}) - \beta(D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}) \| p_\theta(\mathbf{z})) - \delta) \\ &= \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \log p_\theta(\mathbf{x}|\mathbf{z}) - \beta D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}) \| p_\theta(\mathbf{z})) + \beta\delta \\ &\geq \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \log p_\theta(\mathbf{x}|\mathbf{z}) - \beta D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}) \| p_\theta(\mathbf{z})) \quad ; \text{ Because } \beta, \delta \geq 0 \end{aligned}$$

The loss function of β -VAE is defined as:

$$L_{\text{BETA}}(\phi, \beta) = -\mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} \log p_\theta(\mathbf{x}|\mathbf{z}) + \beta D_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}) \| p_\theta(\mathbf{z}))$$

where the Lagrangian multiplier β is considered as a hyperparameter.

Since the negation of $L_{\text{BETA}}(\phi, \beta)$ is the lower bound of the Lagrangian $\mathcal{F}(\theta, \phi, \beta)$. Minimizing the loss is equivalent to maximizing the Lagrangian and thus works for our initial optimization problem.

When $\beta = 1$, it is same as VAE. When $\beta > 1$, it applies a stronger constraint on the latent bottleneck \mathbf{z} and limits the representation capacity of \mathbf{z} . For some conditionally independent generative factors, keeping them disentangled is the most efficient representation. Therefore a high β encourages more

- *VQ loss* : The L2 error between the embedding space and the encoder output
- *Commitment loss*: A measure to encourage the encoder output to stay close to the embedding space and to prevent it from fluctuating too frequently from one code vector to another

$$L = \underbrace{\|\mathbf{x} - D(\mathbf{e}_k)\|_2^2}_{\text{reconstruction loss}} + \underbrace{\|\text{sg}[E(\mathbf{x})] - \mathbf{e}_k\|_2^2}_{\text{VQ loss}} + \underbrace{\beta \|E(\mathbf{x}) - \text{sg}[\mathbf{e}_k]\|_2^2}_{\text{commitment loss}}$$

where `sg[.]` is the `stop_gradient` operator

The embedding vectors in the codebook is updated through EMA (exponential moving average). Given a code vector \mathbf{e}_i , say we have n_i encoder output vectors $\{\mathbf{z}_{i,j}\}_{j=1}^{n_i}$, that are quantized to \mathbf{e}_i :

$$N_i^{(t)} = \gamma N_i^{(t-1)} + (1 - \gamma)n_i^{(t)} \quad \mathbf{m}_i^{(t)} = \gamma \mathbf{m}_i^{(t-1)} + (1 - \gamma) \sum_{j=1}^{n_i^{(t)}} \mathbf{z}_{i,j}^{(t)} \quad \mathbf{e}_i^{(t)} = \mathbf{m}_i^{(t)} / N_i^{(t)}$$

where (t) refers to batch sequence in time, N_i and \mathbf{m}_i are accumulated vector count and volume respectively.

VQ-VAE-2 ([Ali Razavi, et al. 2019](#)) is a two-level hierarchical VQ-VAE combined with self-attention in autoregressive mode

1. Stage 1 is to **train a hierarchical VQ-VAE**: The design of hierarchical latent variables intends to separate local patterns (i.e., texture) from global information (i.e., object shapes). The training of a larger bottom level codebook is conditioned on the smaller top level code book, so that it does not have to learn everything from scratch
2. Stage 2 is to **learn a prior over the latent discrete codebook** so that we sample from it and generate images. In this way, the decoder can receive input vectors sampled from a single distribution as the one in training. A powerful autoregressive model enhanced with multi-head attention layers is used to capture the prior distribution ([PixelSNAIL; Chen et al 2017](#)).

Considering that VQ-VAE-2 depends on discrete latent variables configured in a simple hierarchical setting, the quality of its generated images are pretty amazing




Fig. 11. Architecture of hierarchical VQ-VAE and multi-stage image generation. (Image source: [Ali Razavi, et al. 2019](#))

on the past can be written $ap(x_{t+1}, \dots, x_T | x_1, \dots, x_t) \approx p(x_{t+1}, \dots, x_T | b_t)$. The hidden states in a recurrent policy are used as the agent's belief state in TD-VAE. Thus we have $b_t = \text{RNN}(b_{t-1}, x_t)$.

3. Jumpy Prediction

Further, an agent is expected to imagine distant futures based on all the information gathered: suggesting the capability of making jumpy predictions, that is, predicting states several steps further into the future

Recall what we have learned from the variance lower bound [above](#):

$$\begin{aligned} \log p(x) &\geq \log p(x) - D_{\text{KL}}(q(z|x) \| p(z|x)) \\ &= \mathbb{E}_{z \sim q} \log p(x|z) - D_{\text{KL}}(q(z|x) \| p(z)) \\ &= \mathbb{E}_{z \sim q} \log p(x|z) - \mathbb{E}_{z \sim q} \log \frac{q(z|x)}{p(z)} \\ &= \mathbb{E}_{z \sim q} [\log p(x|z) - \log q(z|x) + \log p(z)] \\ &= \mathbb{E}_{z \sim q} [\log p(x, z) - \log q(z|x)] \\ \log p(x) &\geq \mathbb{E}_{z \sim q} [\log p(x, z) - \log q(z|x)] \end{aligned}$$

Now let's model the distribution of the state x_t as a probability function conditioned on all the past states $x_{<t}$ and two latent variables z_t and z_{t-1} , at current time step and one step back

$$\log p(x_t | x_{<t}) \geq \mathbb{E}_{(z_{t-1}, z_t) \sim q} [\log p(x_t, z_{t-1}, z_t | x_{<t}) - \log q(z_{t-1}, z_t | x_{\leq t})]$$

Continue expanding the equation

$$\begin{aligned} \log p(x_t | x_{<t}) &\geq \mathbb{E}_{(z_{t-1}, z_t) \sim q} [\log p(x_t, z_{t-1}, z_t | x_{<t}) - \log q(z_{t-1}, z_t | x_{\leq t})] \\ &\geq \mathbb{E}_{(z_{t-1}, z_t) \sim q} [\log p(x_t | z_{t-1}, z_t, x_{<t}) + \log p(z_{t-1}, z_t | x_{<t}) - \log q(z_{t-1}, z_t | x_{\leq t})] \\ &\geq \mathbb{E}_{(z_{t-1}, z_t) \sim q} [\log p(x_t | z_t) + \log p(z_{t-1} | x_{<t}) + \log p(z_t | z_{t-1}) - \log q(z_{t-1}, z_t | x_{\leq t})] \\ &\geq \mathbb{E}_{(z_{t-1}, z_t) \sim q} [\log p(x_t | z_t) + \log p(z_{t-1} | x_{<t}) + \log p(z_t | z_{t-1}) - \log q(z_t | x_{\leq t}) - \log q(z_{t-1} | z_t, x_{\leq t})] \end{aligned}$$

Notice two things

- The red terms can be ignored according to Markov assumption
- The blue term is expanded according to Markov assumption
- The green term is expanded to include a one-step prediction back to the past as a smooth distribution

Precisely, there are four types of distributions to learn

- $p_D(\cdot)$ is the **decoder** distribution:
 - $p(x_t | z_t)$ is the encoder by the common definition
 - $p(x_t | z_t) \rightarrow p_D(x_t | z_t)$;
- $p_T(\cdot)$ is the **transition** distribution:
 - $p(z_t | z_{t-1})$ captures the sequential dependency between latent variables
 - $p(z_t | z_{t-1}) \rightarrow p_T(z_t | z_{t-1})$;

3. $p_B(\cdot)$ is the **belief** distribution:

- Both $p(z_{t-1} | x_{<t})$ and $q(z_t | x_{\leq t})$ can use the belief states to predict the latent variable;
- $p(z_{t-1} | x_{<t}) \rightarrow p_B(z_{t-1} | b_{t-1})$;
- $q(z_t | x_{\leq t}) \rightarrow p_B(z_t | b_t)$;

4. $p_S(\cdot)$ is the **smoothing** distribution:

- The back-to-past smoothing term $q(z_{t-1} | z_t, x_{\leq t})$ can be rewritten to be dependent of belief states too;
- $q(z_{t-1} | z_t, x_{\leq t}) \rightarrow p_S(z_{t-1} | z_t, b_{t-1}, b_t)$;

To incorporate the idea of jumpy prediction, the sequential ELBO has to not only work at $t + 1$, but also two distant timestamps $t_1 < t_2$. Here is the final TD-VAE objective function to maximize:

$$J_{t_1, t_2} = \mathbb{E}[\log p_D(x_{t_2} | z_{t_2}) + \log p_B(z_{t_1} | b_{t_1}) + \log p_T(z_{t_2} | z_{t_1}) - \log p_B(z_{t_2} | b_{t_2}) - \log p_S(z_{t_1} | z_{t_2}, b_{t_1}, b_{t_2})]$$




Fig. 14. A detailed overview of TD-VAE architecture, very nicely done. (Image source: [TD-VAE paper](#))

Cited as:

```
@article{weng2018VAE,
    title      = "From Autoencoder to Beta-VAE",
    author     = "Weng, Lilian",
    journal   = "lilianweng.github.io/lil-log",
    year       = "2018",
```

```
url      = "http://lilianweng.github.io/lil-log/2018/08/12/from-autoencoder-to-beta...  
}
```

References

- [1] Geoffrey E. Hinton, and Ruslan R. Salakhutdinov "[Reducing the dimensionality of data with neu networks.](#)" *Science* 313.5786 (2006): 504-507.
- [2] Pascal Vincent, et al. "[Extracting and composing robust features with denoising autoencoc](#)" *ICML*, 2008.
- [3] Pascal Vincent, et al. "[Stacked denoising autoencoders: Learning useful representations in a d](#)" *network with a local denoising criteria* *Journal of machine learning research* 11.Dec (2010) 3371-3408.
- [4] Geoffrey E. Hinton, Nitish Srivastava, Alex Krizhevsky, Ilya Sutskever, and Ruslan R. Salakhutdinov "[Improving neural networks by preventing co-adaptation of feature detectors.](#)" arXiv preprint arXiv:1207.0580 (2012).
- [5] [Sparse Autoencoder](#) by Andrew Ng.
- [6] Alireza Makhzani, Brendan Frey (2013) "[k-sparse autoencoder](#)" *ICLR* 2014.
- [7] Salah Rifai, et al. "[Contractive auto-encoders: Explicit invariance during feature extract](#)" *ICML*, 2011.
- [8] Diederik P. Kingma, and Max Welling "[Auto-encoding variational baye](#)" *ICLR* 2014.
- [9] [Tutorial - What is a variational autoencode](#) on jaan.io
- [10] [Youtube tutorial Variational Autoencoder](#) by Arxiv Insights
- [11] "[A Beginner's Guide to Variational Methods: Mean-Field Approxima](#)" by Eric Jang.
- [12] Carl Doersch. "[Tutorial on variational autoencoder](#)" arXiv:1606.05908, 2016.
- [13] Irina Higgins, et al "[β-VAE: Learning basic visual concepts with a constrained variation framework](#)." *ICLR* 2017.
- [14] Christopher P. Burgess, et al "[Understanding disentangling in beta-VA](#)" *NIPS* 2017.
- [15] Aaron van den Oord, et al "[Neural Discrete Representation Learnin](#)" *NIPS* 2017.
- [16] Ali Razavi, et al. "[Generating Diverse High-Fidelity Images with VQ-VAE](#)." arXiv preprint arXiv:1906.00446 (2019).
- [17] Xi Chen, et al. "[PixelSNAIL: An Improved Autoregressive Generative Model](#)" arXiv preprint arXiv:1712.09763 (2017).

[18] Karol Gregor, et al.“[Temporal Difference Variational Auto-Encoder](#)” ICLR 2019.

← Attention? Attention

Flow-based Deep Generative Models –

2019 © Built by Jekyll and minima | [View this on Github](#) | [Tags](#) | [Contact](#) | [FAQ](#)

